


# “SCALING & AUTOSCALING”



El siguiente documento detallará el paso a paso, el manejo del: **SCALING & AUTOSCALING** de los **CONTENEDORES** en **OPENSHIFT**. La **AGENDA** que se manejará será la siguiente:

- ✓ **REQUERIMIENTOS.**
- ✓ **ESCALAMIENTO DE 'CONTENEDORES'.**
  - **SCALING INSTANCES.**
  - **HORIZONTAL AUTOSCALING INSTANCES.**

DESCRIPCIÓN	DETALLE
<b>1. REQUERIMIENTOS:</b>	
<p>El único requerimiento es tener <b>INSTALADA</b> la herramienta: <b>SIEGE</b> nivel de: <b>VARIABLES DE ENTORNO</b>.</p> <p>Esta herramienta permitirá <b>PROBAR</b> más adelante la aplicación del <b>RECURSO</b>: <b>AUTOSCALING</b>.</p> <ul style="list-style-type: none"><li>▪ <b>WINDOWS:</b></li></ul> <div> siege-windows.zip</div> <pre>\$ siege --version</pre> <ul style="list-style-type: none"><li>▪ <b>LINUX:</b></li></ul> <pre>\$ sudo apt update -y \$ sudo apt install siege -y \$ siege --version</pre> <ul style="list-style-type: none"><li>▪ <b>MAC:</b></li></ul> <pre>\$ brew install siege \$ siege --version</pre>	<pre>GNX-000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master) \$ (siege --version) SIEGE 3.0.5  Copyright (C) 2013 by Jeffrey Fulmer, et al. This is free software; see the source for copying conditions. There is NO warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.</pre>


```
cguerra@LAPTOP-PF8PUGRQ:~$ siege --version
New configuration template added to /home/cguerra/.siege
Run siege -C to view the current settings in that file
SIEGE 4.0.4

Copyright (C) 2017 by Jeffrey Fulmer, et al.
This is free software; see the source for copying conditions.
There is NO warranty; not even for MERCHANTABILITY or FITNESS
FOR A PARTICULAR PURPOSE.
```

```
gianniguatame -- -zsh -- 80x24
gianniguatame@Giannis-MacBook-Pro ~ % [siege --version]
New configuration template added to /Users/gianniguatame/.siege
Run siege -C to view the current settings in that file
SIEGE 4.1.3

Copyright (C) 2022 by Jeffrey Fulmer, et al.
This is free software; see the source for copying conditions.
There is NO warranty; not even for MERCHANTABILITY or FITNESS
FOR A PARTICULAR PURPOSE.
```

## 2. ESCALAMIENTO DE 'CONTENEDORES':



El objetivo aquí es el poder **CRECER** a nivel de **INSTANCIAS (PODS)**, para aplicarlos en los escenarios requeridos. Estos pueden ser: **MANUALMENTE** o **AUTOMÁTICAMENTE**.

### A. SCALING INSTANCES:

Este procedimiento es el **MANUAL** & para ello se debe crear un **ESCENARIO** de prueba, corriendo el **Script YAMEL**:



Script\_Microservicio  
.yaml

```
$ cat > Script_Microservicio.yaml
```

```
...  
.....  
.....
```

```
$ oc create -f Script_Microservicio.yaml
```

```
$ oc get all -n dummy-csm-crga
```

```
GMX+000996815@LAPTOP-PF8PU6RQ MINGW64 ~ (master)
$ cat > Script_Microservicio.yaml
#### ----- [NAMESPACE] ----- ###
apiVersion: v1
kind: Namespace
metadata:
  name: dummy-csm-crga
---
#### ----- [DEPLOYMENT] ----- ###
apiVersion: apps/v1
kind: Deployment
metadata:
  name: dummy-micro-deploy
  namespace: dummy-csm-crga
  labels:
    app: dummy-micro-service
    version: v1
spec:
  replicas: 1
  selector:
    matchLabels:
      app: dummy-micro-service
      version: v1
  template:
    metadata:
      labels:
        app: dummy-micro-service
        version: v1
    spec:
      containers:
        - image: maktup/dummy-micro-01:latest
          name: dummy-micro-container
          resources:
            limits:
              cpu: 250m
              memory: 250Mi
            requests:
              cpu: 100m
              memory: 100Mi
          ports:
            - containerPort: 8080
---
#### ----- [SERVICE] ----- ###
```

	<pre>GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master) \$ oc create -f Script_Microservicio.yaml namespace/dummy-csm-crga created deployment.apps/dummy-micro-deploy created service/dummy-micro-service created route.route.openshift.io/dummy-micro-route created  GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master) \$ oc get all -n dummy-csm-crga</pre> <table><tr><th>NAME</th><th>READY</th><th>STATUS</th><th>RESTARTS</th><th>AGE</th></tr><tr><td>pod/dummy-micro-deploy-c54c7cf69-xs95l</td><td>1/1</td><td>Running</td><td>0</td><td>10s</td></tr></table> <table><tr><th>NAME</th><th>TYPE</th><th>CLUSTER-IP</th><th>EXTERNAL-IP</th><th>PORT(S)</th><th>AGE</th></tr><tr><td>service/dummy-micro-service</td><td>ClusterIP</td><td>172.21.30.139</td><td>&lt;none&gt;</td><td>8080/TCP</td><td>10s</td></tr></table> <table><tr><th>NAME</th><th>READY</th><th>UP-TO-DATE</th><th>AVAILABLE</th><th>AGE</th></tr><tr><td>deployment.apps/dummy-micro-deploy</td><td>1/1</td><td>1</td><td>1</td><td>10s</td></tr></table> <table><tr><th>NAME</th><th>DESIRED</th><th>CURRENT</th><th>READY</th><th>AGE</th></tr><tr><td>replicaset.apps/dummy-micro-deploy-c54c7cf69</td><td>1</td><td>1</td><td>1</td><td>10s</td></tr></table> <table><tr><th>NAME</th><th>PATH</th><th>SERVICES</th><th>PORT</th><th>HOST/PORT</th><th>TERMINATION</th><th>WILDCARD</th></tr><tr><td>route.route.openshift.io/dummy-micro-route</td><td></td><td></td><td></td><td>dummy-micro-route-dummy-csm-crga.cluster-cla-crga-</td><td></td><td></td></tr><tr><td>loud</td><td></td><td>dummy-micro-service</td><td>&lt;all&gt;</td><td></td><td></td><td>None</td></tr></table>	NAME	READY	STATUS	RESTARTS	AGE	pod/dummy-micro-deploy-c54c7cf69-xs95l	1/1	Running	0	10s	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	service/dummy-micro-service	ClusterIP	172.21.30.139	<none>	8080/TCP	10s	NAME	READY	UP-TO-DATE	AVAILABLE	AGE	deployment.apps/dummy-micro-deploy	1/1	1	1	10s	NAME	DESIRED	CURRENT	READY	AGE	replicaset.apps/dummy-micro-deploy-c54c7cf69	1	1	1	10s	NAME	PATH	SERVICES	PORT	HOST/PORT	TERMINATION	WILDCARD	route.route.openshift.io/dummy-micro-route				dummy-micro-route-dummy-csm-crga.cluster-cla-crga-			loud		dummy-micro-service	<all>			None
NAME	READY	STATUS	RESTARTS	AGE																																																												
pod/dummy-micro-deploy-c54c7cf69-xs95l	1/1	Running	0	10s																																																												
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE																																																											
service/dummy-micro-service	ClusterIP	172.21.30.139	<none>	8080/TCP	10s																																																											
NAME	READY	UP-TO-DATE	AVAILABLE	AGE																																																												
deployment.apps/dummy-micro-deploy	1/1	1	1	10s																																																												
NAME	DESIRED	CURRENT	READY	AGE																																																												
replicaset.apps/dummy-micro-deploy-c54c7cf69	1	1	1	10s																																																												
NAME	PATH	SERVICES	PORT	HOST/PORT	TERMINATION	WILDCARD																																																										
route.route.openshift.io/dummy-micro-route				dummy-micro-route-dummy-csm-crga.cluster-cla-crga-																																																												
loud		dummy-micro-service	<all>			None																																																										
<p>Luego, se ingresa el <b>COMANDO</b> para el <b>AUMENTO MANUAL</b> hacia <b>5 INSTANCIAS</b>:</p> <pre>\$ oc scale deployment dummy-micro-deploy --replicas=5 -n dummy-csm-crga</pre> <pre>\$ oc get pods -n dummy-csm-crga</pre>	<pre>GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master) \$ oc scale deployment dummy-micro-deploy --replicas=5 -n dummy-csm-crga deployment.apps/dummy-micro-deploy scaled  GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master) \$ oc get pods -n dummy-csm-crga</pre> <table><tr><th>NAME</th><th>READY</th><th>STATUS</th><th>RESTARTS</th><th>AGE</th></tr><tr><td>dummy-micro-deploy-c54c7cf69-2csf1</td><td>1/1</td><td>Running</td><td>0</td><td>11s</td></tr><tr><td>dummy-micro-deploy-c54c7cf69-h2rdr</td><td>1/1</td><td>Running</td><td>0</td><td>11s</td></tr><tr><td>dummy-micro-deploy-c54c7cf69-phrtj</td><td>1/1</td><td>Running</td><td>0</td><td>11s</td></tr><tr><td>dummy-micro-deploy-c54c7cf69-rpmw8</td><td>1/1</td><td>Running</td><td>0</td><td>11s</td></tr><tr><td>dummy-micro-deploy-c54c7cf69-xs95l</td><td>1/1</td><td>Running</td><td>0</td><td>11m</td></tr></table>	NAME	READY	STATUS	RESTARTS	AGE	dummy-micro-deploy-c54c7cf69-2csf1	1/1	Running	0	11s	dummy-micro-deploy-c54c7cf69-h2rdr	1/1	Running	0	11s	dummy-micro-deploy-c54c7cf69-phrtj	1/1	Running	0	11s	dummy-micro-deploy-c54c7cf69-rpmw8	1/1	Running	0	11s	dummy-micro-deploy-c54c7cf69-xs95l	1/1	Running	0	11m																																	
NAME	READY	STATUS	RESTARTS	AGE																																																												
dummy-micro-deploy-c54c7cf69-2csf1	1/1	Running	0	11s																																																												
dummy-micro-deploy-c54c7cf69-h2rdr	1/1	Running	0	11s																																																												
dummy-micro-deploy-c54c7cf69-phrtj	1/1	Running	0	11s																																																												
dummy-micro-deploy-c54c7cf69-rpmw8	1/1	Running	0	11s																																																												
dummy-micro-deploy-c54c7cf69-xs95l	1/1	Running	0	11m																																																												
<p>Luego, se realiza una prueba del <b>MICROSERVICIO</b> desplegado. Los <b>REQUEST</b> enviados serán distribuidos a nivel de <b>TODAS</b> las <b>INSTANCIAS</b> desplegadas:</p> <pre>\$ curl http://dummy-micro-route-dummy-csm-crga.cluster-cla-crga-ccc03eca20d26e6ac64511f874a64b9b-0000.br-sao.containers.appdomain.cloud/dummy-micro-01/get/personas</pre>	<pre>GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master) \$ curl http://dummy-micro-route-dummy-csm-crga.cluster-cla-crga-ccc03eca20d26e6ac64511f874a64b9b-0000.br-sao.co ntainers.appdomain.cloud/dummy-micro-01/get/personas</pre> <table><tr><th>% Total</th><th>% Received</th><th>% Xferd</th><th>Average Speed</th><th>Time</th><th>Time</th><th>Time</th><th>Current</th></tr><tr><th></th><th></th><th></th><th>Dload</th><th>Upload</th><th>Total</th><th>Spent</th><th>Left</th></tr><tr><td>100</td><td>196</td><td>100</td><td>196</td><td>0</td><td>0</td><td>58</td><td>0</td></tr><tr><td>0:00:03</td><td>0:00:03</td><td>--:--:--</td><td>60</td><td>[{ 'nombre': 'PAOLO GUERRERO', 'e</td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td><td>dad': 35, 'rol': 'CONSULTOR' }, { 'nombre': 'LUIS GUADALUPE', 'edad': 40, 'rol': 'PROGRAMADOR' }, { 'nombre': 'PEDRO SALAZAR', 'edad': 30, 'rol': 'ARQUITECTO' } ]</td><td></td><td></td><td></td></tr></table>	% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current				Dload	Upload	Total	Spent	Left	100	196	100	196	0	0	58	0	0:00:03	0:00:03	--:--:--	60	[{ 'nombre': 'PAOLO GUERRERO', 'e								dad': 35, 'rol': 'CONSULTOR' }, { 'nombre': 'LUIS GUADALUPE', 'edad': 40, 'rol': 'PROGRAMADOR' }, { 'nombre': 'PEDRO SALAZAR', 'edad': 30, 'rol': 'ARQUITECTO' } ]																										
% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current																																																									
			Dload	Upload	Total	Spent	Left																																																									
100	196	100	196	0	0	58	0																																																									
0:00:03	0:00:03	--:--:--	60	[{ 'nombre': 'PAOLO GUERRERO', 'e																																																												
				dad': 35, 'rol': 'CONSULTOR' }, { 'nombre': 'LUIS GUADALUPE', 'edad': 40, 'rol': 'PROGRAMADOR' }, { 'nombre': 'PEDRO SALAZAR', 'edad': 30, 'rol': 'ARQUITECTO' } ]																																																												

Finalmente, para el **DECREMENTO MANUAL** de **INSTANCIAS** hacia: **1** (dejarlo como estaba inicialmente), se ingresa:

```
$ oc scale deployment dummy-micro-deploy --replicas=1 -n dummy-csm-crga
```

```
$ oc get pods -n dummy-csm-crga
```

```
GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc scale deployment dummy-micro-deploy --replicas=1 -n dummy-csm-crga
deployment.apps/dummy-micro-deploy scaled

GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc get pods -n dummy-csm-crga
```

NAME	READY	STATUS	RESTARTS	AGE
dummy-micro-deploy-c54c7cf69-2csf1	0/1	Terminating	0	10m
dummy-micro-deploy-c54c7cf69-h2rdr	1/1	Running	0	10m
dummy-micro-deploy-c54c7cf69-phrtj	0/1	Terminating	0	10m
dummy-micro-deploy-c54c7cf69-rpmw8	0/1	Terminating	0	10m
dummy-micro-deploy-c54c7cf69-xs95l	0/1	Terminating	0	21m

```
GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc get pods -n dummy-csm-crga
```

NAME	READY	STATUS	RESTARTS	AGE
dummy-micro-deploy-c54c7cf69-h2rdr	1/1	Running	0	13m

## B. HORIZONTAL AUTOSCALING INSTANCES:

El objetivo aquí es el poder **CRECER** a nivel de **INSTANCIAS (PODs)**, esto aplicado en los escenarios de **TRÁFICO CAMBIANTE**, sobre todo, pero de manera: **AUTOMÁTICA**.

**IMPORTANTE:** “Antes que nada, se debe conocer que el **HORIZONTAL AUTOSCALING**, requiere que esté **ACTIVO** algo que se llama: **METRIC-SERVER**, esto sirve para poder desde **OPENSIFT** medir el control **CPU & MEMORIA** tanto a nivel de: **NODOS & PODS** respectivamente.

En **KUBERNETES** no está instalado, pero si en **OPENSIFT (NO se requiere)**”:

```
$ git clone https://github.com/kodekloudhub/kubernetes-metrics-server.git
$ cd kubernetes-metrics-server
$ oc create -f ./
```

Luego, para **PROBAR** este escenario se requiere ejecutar los **COMANDOS IMPERATIVOS** siguientes:

```
$ oc autoscale deployment dummy-micro-deploy --min=1 --max=4 --cpu-percent=40 -n dummy-csm-crga
```

```
$ oc get pods,hpa -o wide -n dummy-csm-crga
```

**IMPORTANTE:** “Esta modalidad a nivel de **COMANDOS**, aplica una **REGLA** que funciona solo a nivel de **CPU**”.

```
GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc autoscale deployment dummy-micro-deploy --min=1 --max=4 --cpu-percent=40 -n dummy-csm-crga
horizontalpodautoscaler.autoscaling/dummy-micro-deploy autoscaled

GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc get pods,hpa -o wide -n dummy-csm-crga
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINA
pod/dummy-micro-deploy-c54c7cf69-h2rdr	1/1	Running	0	26m	172.30.240.218	10.150.181.189	<none>

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS
horizontalpodautoscaler.autoscaling/dummy-micro-deploy	Deployment/dummy-micro-deploy	3%/40%	1	4

Luego, es necesario obtener el **ROUTE** creado del **MICROSERVICIO**, ejecutando:

```
$ oc get routes -n dummy-csm-crga
```

```
GNX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc get routes -n dummy-csm-crga
NAME                                HOST/PORT
---                                -
dummy-micro-route                  dummy-micro-route-dummy-csm-crga.cluster-cla-crga-ccc03eca20d26e6ac64511f874a64b9b-0000.br-sao.containers.appdomain.cloud
dummy-micro-service                <all>
dummy-micro-service                None
```

Luego, se requiere probar **MASIVAMENTE** contra el **MICROSERVICIO** desplegado. Así mismo, los **REQUEST** serán distribuidos a nivel de **TODAS** las **INSTANCIAS** desplegadas y/o se vayan **ACTIVANDO**:

**SIEGE** permitirá enviar **REQUEST** de manera **PARALELA**, esto significa que en la demostración se enviarán: **4 PETICIONES**, cada una con **20 CONCURRENCIAS** (en **TOTAL** se deberán ejecutar: **80 Mensajes**).

```
$ siege --concurrent=20 --reps=4 -v http://dummy-micro-
route-dummy-csm-crga.cluster-cla-crga-
ccc03eca20d26e6ac64511f874a64b9b-0000.us-
south.containers.appdomain.cloud/dummy-micro-
01/get/personas
```

**IMPORTANTE:** “La ejecución del **TRÁFICO MASIVO** hará que el **CPU**, como se muestra la **IMAGEN**, se incremente & genere la **ELASTICIDAD** a nivel de **INSTANCIAS** de **PODs**”.

```
$ oc get pods,hpa -n dummy-csm-crga
```

```
cguerra@LAPTOP-PFBPU6RQ: ~$ oc get pods,hpa -n dummy-csm-crga
NAME                                READY   STATUS    RESTARTS   AGE
pod/dummy-micro-deploy-c54c7cf69-h2dr 1/1     Running   0           136m
pod/dummy-micro-deploy-c54c7cf69-t84k5 1/1     Running   0           11s
```

```
NAME                                REFERENCE
horizontalpodautoscaler.autoscaling/dummy-micro-deploy Deployment/dummy-micro-deploy
```

```
GNX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc get pods,hpa -n dummy-csm-crga
NAME                                READY   STATUS    RESTARTS   AGE
pod/dummy-micro-deploy-c54c7cf69-h2dr 1/1     Running   0           136m
pod/dummy-micro-deploy-c54c7cf69-1xcm  1/1     Running   0           22s
pod/dummy-micro-deploy-c54c7cf69-t84k5 1/1     Running   0           2m24s
```

```
NAME                                REFERENCE
horizontalpodautoscaler.autoscaling/dummy-micro-deploy Deployment/dummy-micro-deploy
```

TARGETS	MINPODS	MAXPODS
5/5/40%	1	4

TARGETS	MINPODS	MAXPODS
110%/40%	1	4

```
cguerra@LAPTOP-PFBPU6RQ: ~$ oc get pods,hpa -n dummy-csm-crga
NAME                                READY   STATUS    RESTARTS   AGE
pod/dummy-micro-deploy-c54c7cf69-h2dr 1/1     Running   0           136m
pod/dummy-micro-deploy-c54c7cf69-1xcm  1/1     Running   0           22s
pod/dummy-micro-deploy-c54c7cf69-t84k5 1/1     Running   0           2m24s
```

```
NAME                                REFERENCE
horizontalpodautoscaler.autoscaling/dummy-micro-deploy Deployment/dummy-micro-deploy
```

```
GNX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc get pods,hpa -n dummy-csm-crga
NAME                                READY   STATUS    RESTARTS   AGE
pod/dummy-micro-deploy-c54c7cf69-h2dr 1/1     Running   0           136m
pod/dummy-micro-deploy-c54c7cf69-1xcm  1/1     Running   0           22s
pod/dummy-micro-deploy-c54c7cf69-t84k5 1/1     Running   0           2m24s
```

```
NAME                                REFERENCE
horizontalpodautoscaler.autoscaling/dummy-micro-deploy Deployment/dummy-micro-deploy
```

(VERSIÓN LINUX)

```
gianniguatame --zsh -- 80x66
gianniguatame@Giannis-MacBook-Pro ~ % siege --concurrent=20 --reps=4 -v http://
dummy-micro-route-dummy-csm-crga.cluster-cta-crga-ccc03eca20d26e6ac64511f874a64t
ph-0000.br-sao.containers.apdomain.cloud/dummy-micro-01/get/personas
** SIEGE 4.1.3
** Preparing 20 concurrent users for battle.
The server is now under siege...
HTTP/1.1 200 0.90 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.00 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.10 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.10 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.10 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.20 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.40 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.30 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 1.40 secs: 196 bytes ==> GET /dummy-micro-01/get/personas
```

(VERSIÓN MAC)

Luego, se debe verificar que después de un **TIEMPO** de realizada la **PRUEBA** de: **REQUEST MASIVOS**, la plataforma se debe haber estabilizado a nivel de **CPU**. (La **ELIMINACIÓN** de los **PODs** adicionales es **AUTOMÁTICA**).

```
$ oc get pods,hpa -n dummy-csm-crga
```

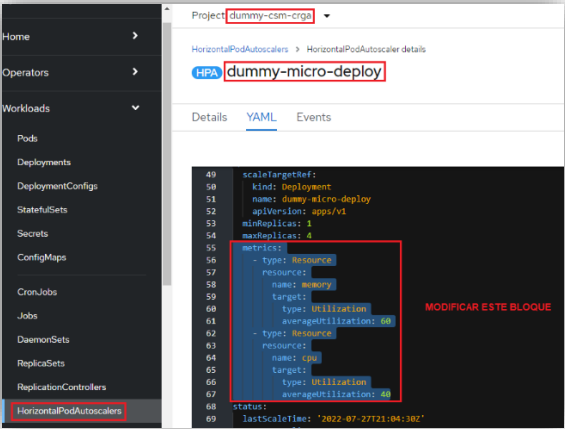
```
GMX+000996815@LAPTOP-PFBPU6RQ MINGW64 ~ (master)
$ oc get pods,hpa -n dummy-csm-crga
NAME                                     READY   STATUS    RESTARTS   AGE
pod/dummy-micro-deploy-c54c7cf69-h2rdr  1/1     Running   0           146m

NAME                                     REFERENCE                                TARGETS   MINPODS   MAXPODS
horizontalpodautoscaler.autoscaling/dummy-micro-deploy  Deployment/dummy-micro-deploy  4%/40%    1         4
```

Finalmente, solo para **COMPLEMENTAR** que la modalidad en **SCRIPT** del: **HORIZONTAL AUTOSCALING**, **NO** solamente funciona a nivel de **CPU**, sino que también funciona **SIMULTÁNEAMENTE** a nivel de **MEMORIA**.

“Esto vuelve más **SENSIBLE** la activación del: **HORIZONTAL AUTOSCALING**”.

```
apiVersion: autoscaling/v2beta2
kind: HorizontalPodAutoscaler
metadata:
  name: dummy-micro-deploy
  namespace: dummy-csm-crga
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: dummy-micro-deploy
  minReplicas: 1
  maxReplicas: 4
  metrics:
  - type: Resource
    resource:
      name: cpu
```



**IMPORTANTE:** “Se aprecia como la **REGLA** del **HORIZONTAL AUTOSCALING**, ahora está aplicara a nivel de: **CPU & MEMORIA**”.

```

$ kubectl exec -i $(kubectl get pods -n dummy -l app=dummy_micro-route-dummy -o jsonpath='{.items[0].metadata.name}') -- curl -s -H 'Host: dummy-micro-route-dummy.csm-crga.cluster-cls-crga-cn3.svc.cluster.local' http://dummy-micro-route-dummy.csm-crga.cluster-cls-crga-cn3.svc.cluster.local:8080/0000.br
sao.containers.appdomain.cloud/dummy-micro-01/get/personas

** SIEGE 3.0.5

The server is now under siege...
done.

Transactions:      80 hits
Availability:      100.00 %
Elapsed time:      4.55 secs
Data transferred:  0.01 MB
Response time:     0.26 secs
Transaction rate:  16.48 trans/sec
Throughput:        0.00 MB/sec
Concurrency:       4.27
Successful Transactions: 80
Failed transactions: 0
Longest transaction: 0.64
Shortest transaction: 0.18

HTTP/1.1 200 0.21 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.21 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.21 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.22 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.22 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.22 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.22 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.26 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.26 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.26 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.26 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.30 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.33 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.33 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.33 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.40 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.40 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas
HTTP/1.1 200 0.64 secs:: 196 bytes ==> GET /dummy-micro-01/get/personas

```